

# The Pilot Of Hadoop

2017.2 XenRon

# **L** CONTENTS

Preliminary Topics 事前準備 01

05 Hadoop Command

Hadoop Eco System

02

Reference Books

Hadoop Architecture **03** 

Hadoop Environment 実験環境

04





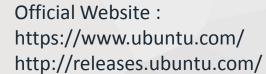
Preliminary Topics 事前準備











Official Website:
https://www.centos.org/
https://www.centos.org/download/
https://wiki.centos.org/Download



Official Website : https://www.redhat.com/en

#### SSH Terminal



4 b

```
69
                                                                                                     sshd
                                                                                                                      0.0
                                                                                                                            00:00.10 2
                                                                                                                                                      37
                                                                                                                                                                    468K
                                                                                                                                                                             3556K
                                                                                                                                                                                       4700K
                                                                                                     usbmuxd
                                                                                                                      0.0 00:01.00 3
                                                                                                                                                      35
                                                                                                                                                             73
                                                                                                                                                                    852K
                                                                                                                                                                             1392K
                                                                                                                                                                                      1892K
                                                                                            526-
                                                                                                                                                             47
                                                                                                                                                                    156K
                                                                                                                                                                             1724K 456K
                                                                                                     thnuclnt
                                                                                                                      0.0 01:17.55 1
                                                                                                                                                0
                                                                                                                                                      8
iTerm2 Shell Edit View Profiles Toolbelt Window Help
                                                                                            525-
                                                                                                    thnuc1nt
                                                                                                                      0.0 00:05.86 1
                                                                                                                                                0
                                                                                                                                                      8
                                                                                                                                                             47
                                                                                                                                                                    408K
                                                                                                                                                                             1724K
                                                                                                                                                                                       568K
                                                                                            519-
                                                                                                     thnuclnt
                                                                                                                      0.0 00:00.00 1
                                                                                                                                                      9
                                                                                                                                                             47
                                                                                                                                                                    28K
                                                                                                                                                                             1724K
                                                                                                                                                                                       308K
                                                                                           518-
                                                                                                     thnuclnt
                                                                                                                      0.0
                                                                                                                            00:00.00 1
                                                                                                                                                0
                                                                                                                                                      8
                                                                                                                                                             45
                                                                                                                                                                    28K
                                                                                                                                                                             1724K
                                                                                                                                                                                       308K
                                                                                           517-
                                                                                                    thnuclnt
                                                                                                                      0.0 01:24.92 2
                                                                                                                                                      21
                                                                                                                                                             63
                                                                                                                                                                    3468K
                                                                                                                                                                             1680K
                                                                                                                                                                                       6884K
          ExprNode *root = buildAssignment( iter, vars, funs );
                                                                                           515-
                                                                                                    thnuc1nt
                                                                                                                      0.0 00:06.82 1
                                                                                                                                                      17
                                                                                                                                                             45
                                                                                                                                                                    812K
                                                                                                                                                                             1720K
                                                                                                                                                                                      1624K
          return root->evaluate( vars );
                                                                                                                                                             611
                                                                                                     vmware-vmx
                                                                                                                      2.8 02:57:02 24
                                                                                                                                                      250
                                                                                                                                                                    31 M
                                                                                                                                                                             53M
                                                                                                                                                                                       1343M
                                                                                    ,#####:
                                                                                                                    ssh2: AES-256-CTR
                                                                                                                                          24. 74 24 Rows, 80 Cols Linux
                                                                                                                                                                                  00:01:18 CAP NUM
                                                                                  ########
                                                                                  #########
                                                                                  Terminal: xterm-256color
                                                                                  #############################
                                                                                                         Packages: 52
                                                                                                         CPU: Intel Core i5-4258U CPU @ 2.40GHz
                                                                                                         Memory: 8 GB
 66 ExprNode* buildAssignment( ListIterator &iter, VarTree &vars, FunctionDef &funs )
                                                                                   Graphics: Iris
                                                                                    #################
                                                                                                         Disk: 202G / 250G. 82%
                                                                                     '####'''####'
                                                                                                         IP: 104.38.35.203
      if (iter.tokenChar() != '(' && !iter.currentIsInteger())
          ListIterator temp = iter; // make a copy
                                                                              client-104-38-35-203:0SXey drew$ |
          temp.advance(); // move forward to check for an '='
          if (!temp.token().isNull())
                                                                              / The eternal feminine draws us upward. \
             if (temp.token().tokenText() == "=") // there is an assignment taking place
                                                                              \ -- Goethe
                ExprNode* left = new Variable( iter.token().tokenText() );
                iter.advance(); //advance to '='
                iter.advance(): // advance past
```

PID

20629

20624

20620

0619

0618

peterbilt.bigrig.org - SecureCRT

COMMAND

mdworker

mdworker

top

bash

sshd

File Edit View Options Transfer Script Tools Help

Disks: 412593/16G read, 249298/13G written.

0.0

0.0

SharedLibs: 5744K resident, 556K data, OB linkedit.

Networks: packets: 4743894/901m in, 4664408/603m out.

3.4 00:01.36 1/1

0.0 00:00.06 3

0.0 00:00.01 1

00:00.06 3

00:00.01 1

%CPU TIME

Processes: 75 total, 3 running, 1 stuck, 71 sleeping, 285 threads

4emRegions: 8243 total. 301M resident. 17M private. 1257M shared.

oad Avg: 0.25, 0.32, 0.24 CPU usage: 0.97% user, 3.88% sys, 95.14% idle

PhysMem: 1555M wired, 1250M active, 1283M inactive, 4087M used, 10M free. vM: 157G vsize, 1033M framework vsize, 407567(0) pageins, <u>33679(0) pageouts.</u>

#TH #WO

0

0

**#POR #MRE RPRVT** 

944K

324K

1600K

268K

1512K

44

61

25

62

69

24

48

48

11

RSHRD

264K

244K

3556K

14M

10M

RSIZE

1516K

4236K

4532K

996K

780K

🛩 peterbilt.bigrig.org 🛛 🗙 💚 mack.bigrig.org









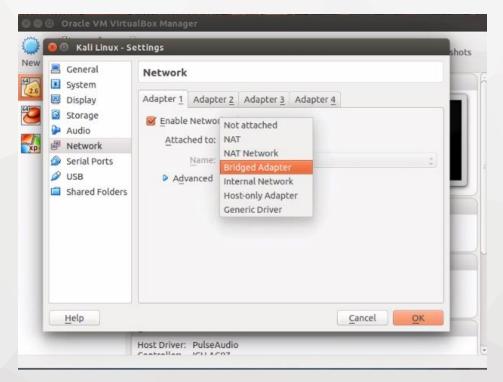












Official Website:

https://www.virtualbox.org/

https://www.virtualbox.org/wiki/Downloads

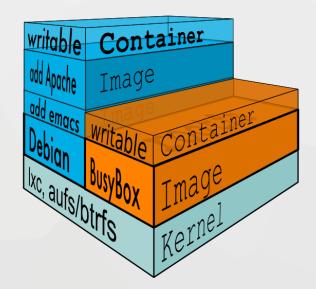
Virtual networking:

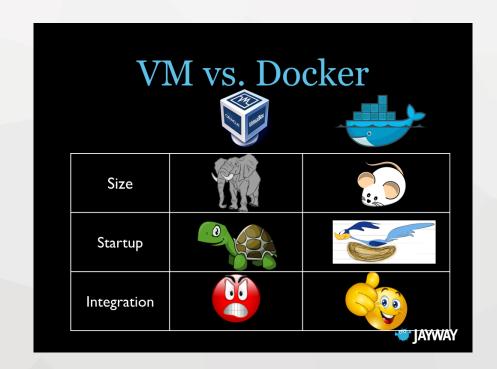
https://www.virtualbox.org/manual/ch06.html







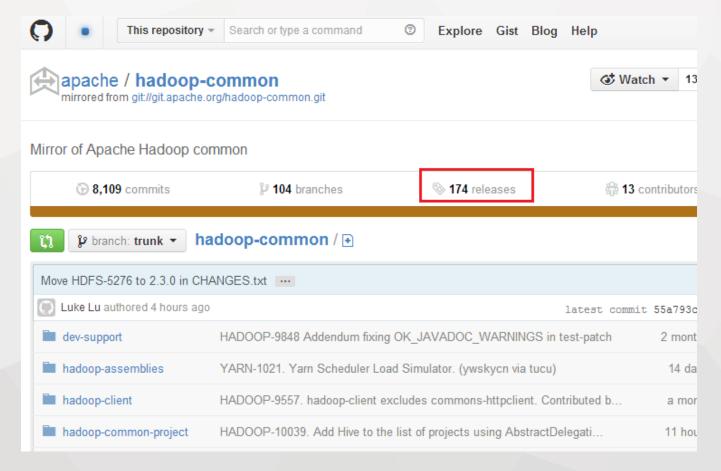








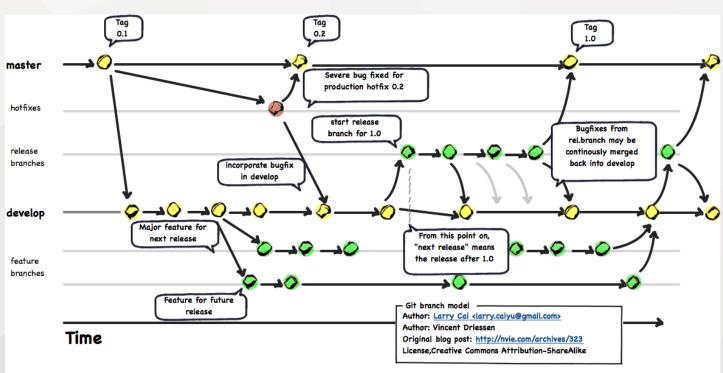












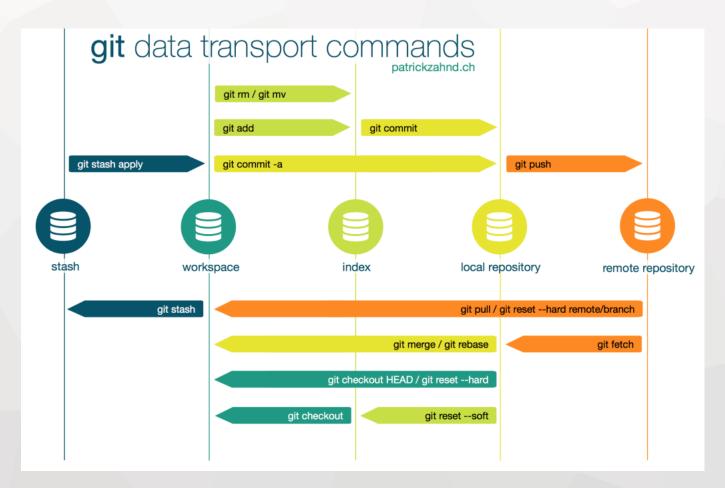
Official Website:

https://git-scm.com/

The latest stable Release: v2.11.1





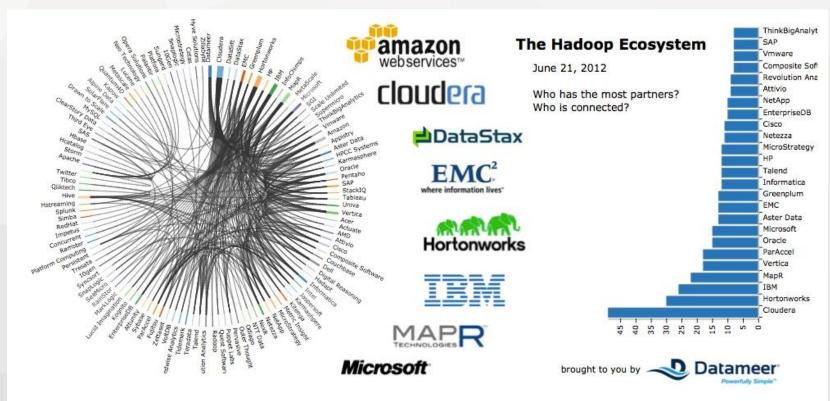




**Hadoop Eco System** 

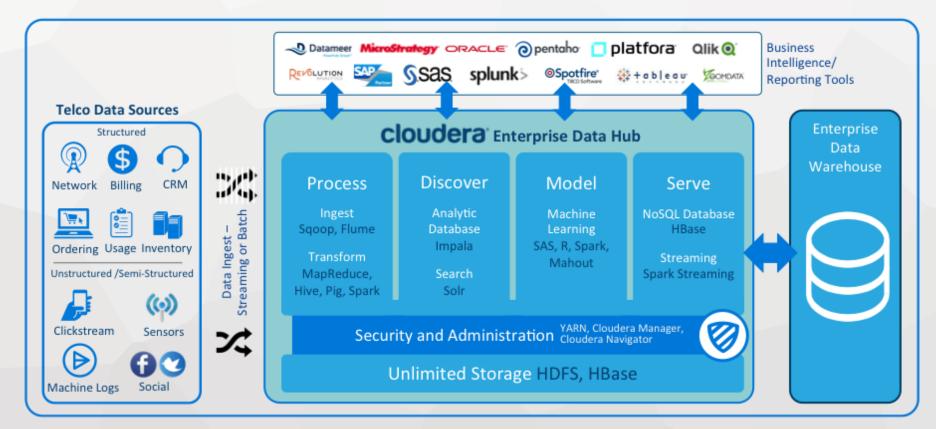
#### Distributions





#### Cloudera





### | Hadoop Eco System

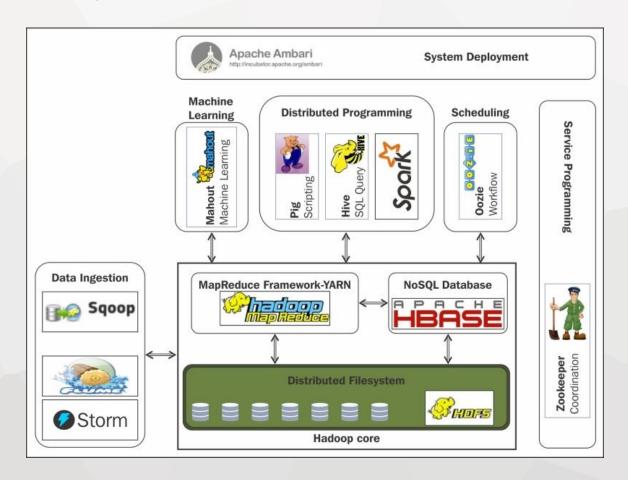






### Hadoop Eco System



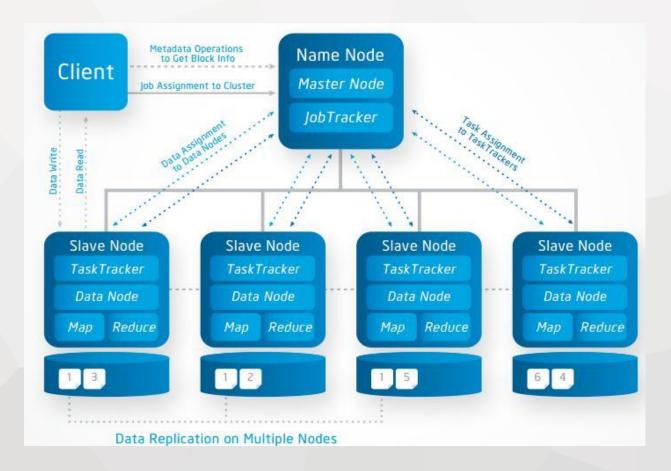




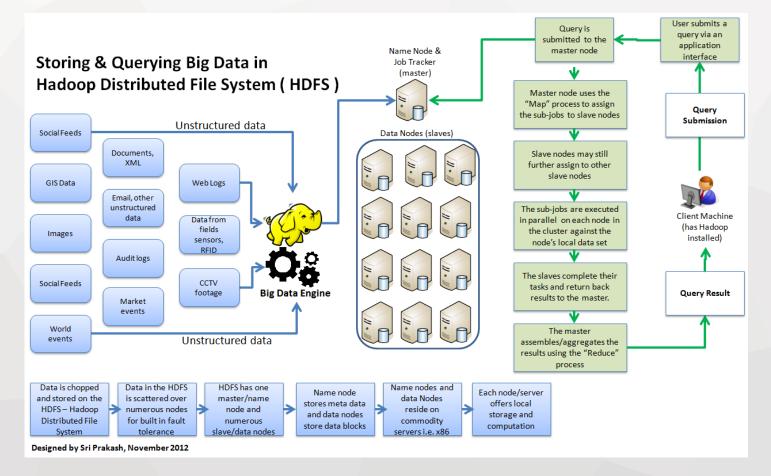
**Hadoop Architecture** 

# Architecture



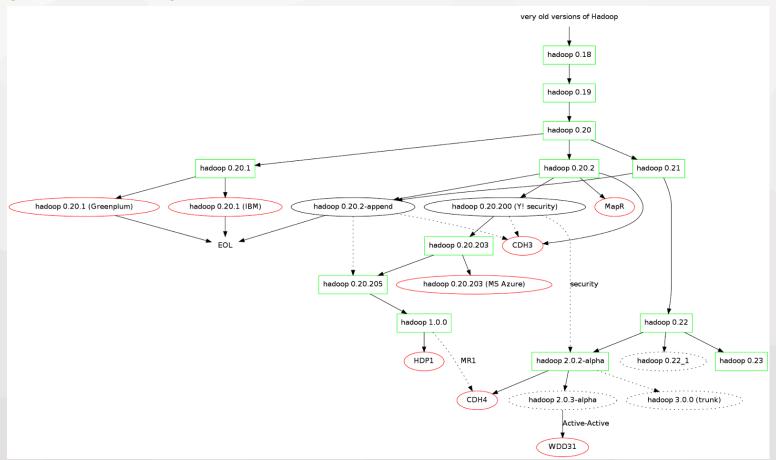






# Apache Hadoop Version



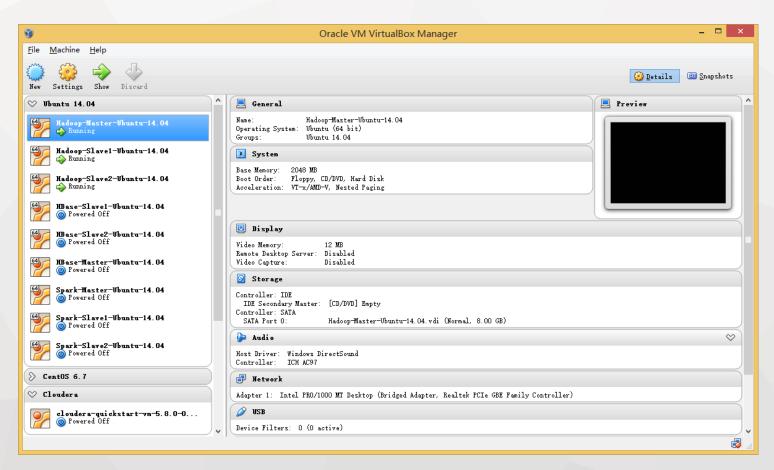


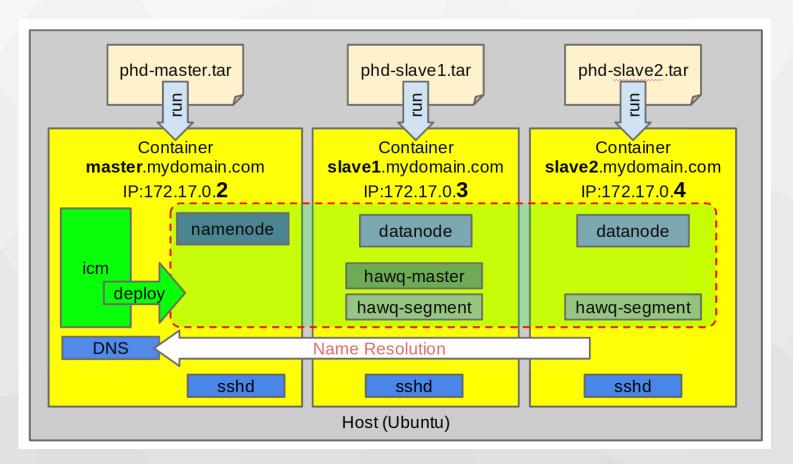


**Hadoop Environment** 

#### **Oracle Virtual Box**















## Hadoop configuration files



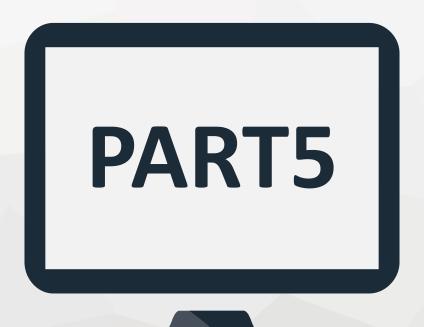
Table 9-1. Hadoop configuration files

Filename	Format	Description
hadoop-env.sh	Bash script	Environment variables that are used in the scripts to run Hadoop
core-site.xml	Hadoop configuration XML	Configuration settings for Hadoop Core, such as I/O settings that are common to HDFS and MapReduce
hdfs-site.xml	Hadoop configuration XML	Configuration settings for HDFS daemons: the namenode, the secondary namenode, and the datanodes
mapred-site.xml	Hadoop configuration XML	Configuration settings for MapReduce daemons: the jobtracker, and the tasktrackers
masters	Plain text	A list of machines (one per line) that each run a secondary namenode
slaves	Plain text	A list of machines (one per line) that each run a datanode and a task-tracker
hadoop-metrics .prop- erties	Java Properties	Properties for controlling how metrics are published in Hadoop (see "Metrics" on page 352)
log4j.properties	Java Properties	Properties for system logfiles, the namenode audit log, and the task log for the tasktracker child process ("Hadoop Logs" on page 175)

Hadoop 1.x Hadoop.The.Definitive.Guide.3rd.Edition

	Tab	le 10-1. Hadoop configuration files
Filename	Format	Description
hadoop-env.sh	Bash script	Environment variables that are used in the scripts to run Hadoop
mapred-env.sh	Bash script	Environment variables that are used in the scripts to run MapReduce (overrides variables set in <i>hadoop-env.sh</i> )
yarn-env.sh	Bash script	Environment variables that are used in the scripts to run YARN (overrides variables set in <i>hadoop-env.sh</i> )
core-site.xml	Hadoop configuration XML	Configuration settings for Hadoop Core, such as I/O settings that are common to HDFS, MapReduce, and YARN
hdfs-site.xml	Hadoop configuration XML	Configuration settings for HDFS daemons: the namenode, the secondary namenode, and the datanodes
mapred-site.xml	Hadoop configuration XML	Configuration settings for MapReduce daemons: the job history server
yarn-site.xml	Hadoop configuration XML	Configuration settings for YARN daemons: the resource manager, the web app proxy server, and the node managers
slaves	Plain text	A list of machines (one per line) that each run a datanode and a node manager
hadoop-metrics2 .properties	Java properties	Properties for controlling how metrics are published in Hadoop (see Metrics and JMX)
log4j.properties	Java properties	Properties for system logfiles, the namenode audit log, and the task log for the task JVM process (Hadoop Logs)
hadoop-policy.xml	Hadoop configuration XML	Configuration settings for access control lists when running Hadoop in secure mode

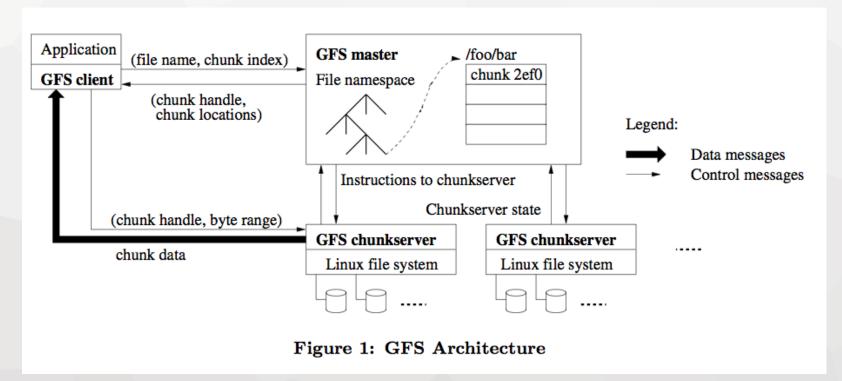
Hadoop 2.x Hadoop.The.Definitive.Guide.4th.Edition



**HDFS & Hadoop Command** 

#### Google File System





The Google File System https://static.googleusercontent.com/media/research.google.com/ja//archive/gfs-sosp2003.pdf



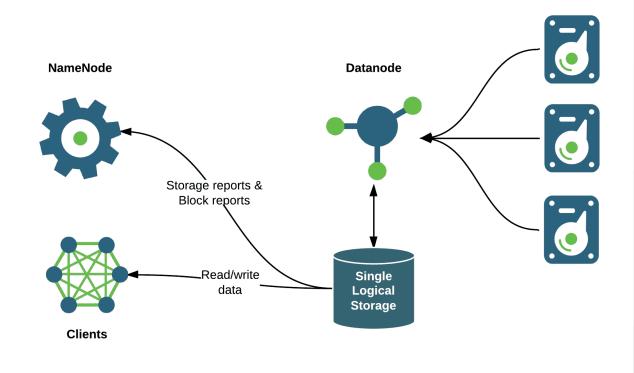


Figure 1: A DataNode presented itself as a single logical storage





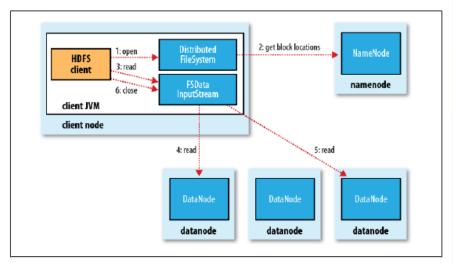


Figure 3-2. A client reading data from HDFS

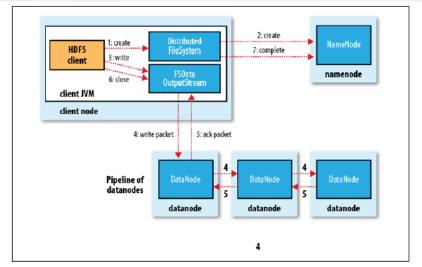


Figure 3-4. A client writing data to HDFS





```
🕒 🗈 hduser@kannandreams: /usr/local/hadoop/bin
Usage: hadoop [--config confdir] COMMAND
where COMMAND is one of:
  namenode -format
                       format the DFS filesystem
  secondarynamenode
                       run the DFS secondary namenode
  namenode
                       run the DFS namenode
  datanode
                       run a DFS datanode
  dfsadmin
                       run a DFS admin client
  mradmin
                       run a Map-Reduce admin client
                       run a DFS filesystem checking utility
  fsck
                       run a generic filesystem user client
  balancer
                       run a cluster balancing utility
                       apply the offline fsimage viewer to an fsimage
  oiv
                       fetch a delegation token from the NameNode
  fetchdt
  iobtracker
                       run the MapReduce job Tracker node
                       run a Pipes job
  pipes
                       run a MapReduce task Tracker node
  tasktracker
  historyserver
                       run job history servers as a standalone daemon
                       manipulate MapReduce jobs
  iob
                       get information regarding JobQueues
  queue
                       print the version
  version
                       run a jar file
  iar <iar>
  distcp <srcurl> <desturl> copy file or directories recursively
  distcp2 <srcurl> <desturl> DistCp version 2
  archive -archiveName NAME -p <parent path> <src>* <dest> create a hadoop archive
  classpath
                       prints the class path needed to get the
                       Hadoop jar and the required libraries
  daemonlog
                       get/set the log level for each daemon
 οг
                       run the class named CLASSNAME
  CLASSNAME
Most commands print help when invoked w/o parameters.
hduser@kannandreams:/usr/local/hadoop/bin$
```

#### Official Website:

https://hadoop.apache.org/docs/s table/hadoop-projectdist/hadoopcommon/FileSystemShell.html

https://hadoop.apache.org/docs/r 1.0.4/cn/hdfs\_shell.html





-	-		_	_	_	_
			_	0 (	_	_
	_					_
		_	-	•	•	_

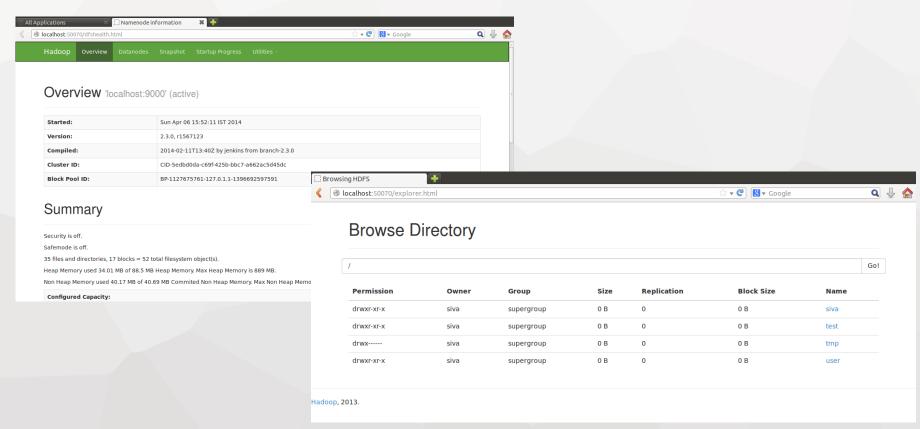


export JAVA_HOME=/usr/java/jdk1.6.0_26/ for service in /etc/init.d/hadoop-0.20-* do sudo \$service start done	STARTING THE PROCESSES
Hadoop NameNode (http:// <hostname>:50070) Hadoop JobTracker (http://<hostname>:50030)</hostname></hostname>	MONITORING PAGES
for service in /etc/init.d/hadoop-0.20-* do sudo \$service stop done	STOPPING THE PROCESSES

HADOOP HDFS			
hadoop fs -ls <path></path>	LIST FILES		
hadoop fs -mkdir <path></path>	MAKE DIRECTORY		
hadoop fs -rmr <path></path>	REMOVE DIRECTORY		
hadoop fs -put <local_file> <hdfs_path></hdfs_path></local_file>	LOAD FILE		
hadoop fs -rm <file></file>	REMOVE FILE		
hadoop fs -cat <file> hadoop fs -tail <file></file></file>	VIEW FILE		
hadoop fs -getmerge <hdfs_ directory&gt; <local_output_file></local_output_file></hdfs_ 	MERGE MULTIPLE PART FILES		

# HDFS Web UI







**Reference Books** 



